



# Local Search

- Tree search keeps unexplored alternatives on the fringe (ensures completeness)
- Local search: improve a single option until you can't make it better (no fringe!)
- New successor function: local changes
- Generally much faster and more memory efficient (but incomplete and suboptimal)




# Hill Climbing


- Simple, general idea:
  - Start wherever
  - Repeat: move to the best neighboring state
  - If no neighbors better than current, quit
- What's bad about this approach?
  - Complete?
  - Optimal?
- What's good about it?



# Hill Climbing Diagram



# Hill Climbing Quiz



Starting from  $X$ , where do you end up ?

Starting from  $Y$ , where do you end up ?

Starting from  $Z$ , where do you end up ?

# Hill climbing search

- Here one always chooses a successor  $s' \in S(s)$  of the current state  $s$  that has the highest value for the objective function  $f$

$$\max_{s' \in S(s)} f(s')$$

- Search terminates when all neighbors of the state have a lower value for the objective function than the current state has
- Most often search terminates in a local maximum, sometimes by chance, in a global maximum
- Also plateaux cause problems to this greedy local search
- On the other hand, improvement starting from the initial state is often very fast


- Sideways moves can be allowed when search may proceed to states that are as good as the current one
- **Stochastic** hill-climbing chooses at random one of the neighbors that improve the situation
- Neighbors can, for example, be examined in random order and choose the first one that is better than the current state
- Also these versions of hill-climbing are incomplete because they can still get stuck in a local maximum
- By using **random restarts** one can guarantee the completeness of the method
  - Start from a random initial state until a solution is found

# Local beam search


- The search begins with  $k$  randomly generated states
- At each step, all the successors of all  $k$  states are generated
- If any one of the successors is a goal, the algorithm halts
- Otherwise, it selects the  $k$  best successors from the complete list and repeats
- The parallel search leads quickly to abandoning unfruitful searches and moves its resources to where the most progress is being made
- In stochastic beam search the maintained successor states are chosen with a probability based on their goodness

# Beam Search

- Like greedy hill climbing search, but keep  $K$  states at all times:



Greedy Search




Beam Search

- Variables: beam size, encourage diversity?
- The best choice in MANY practical settings
- Complete? Optimal?
- Why do we still need optimal methods?



# Random walk

- A **random walk** moves to a successor chosen uniformly at random independent of whether it is better than the current state
  - a complete search algorithm, but when unsupervised also extremely inefficient
- Let us allow “bad” moves with some probability  $p$
- The probability of transitions leading to worse situation decreases exponentially with time (“temperature”)
- We choose a candidate for transition randomly and accept it if
  - the objective value improves;
  - otherwise with probability  $p$
- If temperature is lowered slowly enough, this method converges to a global optimum with probability  $\rightarrow 1$




# Simulated Annealing

- Idea: Escape local maxima by allowing downhill moves
  - But make them rarer as time goes on

```

function SIMULATED-ANNEALING(problem, schedule) returns a solution state
    inputs: problem, a problem
              schedule, a mapping from time to “temperature”
    local variables: current, a node
                      next, a node
                      T, a “temperature” controlling prob. of downward steps

    current ← MAKE-NODE(INITIAL-STATE[problem])
    for t ← 1 to ∞ do
        T ← schedule[t]
        if T = 0 then return current
        next ← a randomly selected successor of current
         $\Delta E \leftarrow \text{VALUE}[\textit{next}] - \text{VALUE}[\textit{current}]$ 
        if  $\Delta E > 0$  then current ← next
        else current ← next only with probability  $e^{\Delta E/T}$ 
    
```




# Simulated Annealing

- Theoretical guarantee:



- Stationary distribution:

$$p(x) \propto e^{\frac{E(x)}{kT}}$$

- If  $T$  decreased slowly enough,  
will converge to optimal state!
- Is this an interesting guarantee?
- Sounds like magic, but reality is reality:
  - The more downhill steps you need to escape a local optimum, the less likely you are to ever make them all in a row
  - People think hard about *ridge operators* which let you jump around the space in better ways




# Genetic Algorithms



- Genetic algorithms use a natural selection metaphor
  - Keep best  $N$  hypotheses at each step (selection) based on a fitness function
  - Also have pairwise crossover operators, with optional mutation to give variety
- Possibly the most misunderstood, misapplied (and even maligned) technique around

# Example: $N$ -Queens



- Why does crossover make sense here?
- When wouldn't it make sense?
- What would mutation be?
- What would a good fitness function be?

# Online Search

- Not all offline search algorithms are suitable for online search
- E.g.,  $A^*$  is essentially based on the fact that one can expand any node generated to the search tree
- An online algorithm can expand only a node that it physically occupies
- DFS only uses local information, except when backtracking
- Hence, it is usable in online search (if actions can physically be undone)
- DFS is not competitive: one cannot bound the competitive ratio



- Hill-climbing search already is an online algorithm, but it gets stuck at local maxima
- Random restarts cannot be used:
  - the agent cannot transport itself to a new state
- Random walks are too inefficient
- Using extra space may make hill-climbing useful in online search
- We store for each state  $s$  visited our current best estimate  $h(s)$  of the cost to reach the goal
- Rather than staying where it is, the agent follows what seems to be the best path to the goal based on the current  $h(s)$  cost estimates for its neighbors
- At the same time the value of a local minimum gets flattened out and can be escaped

