


Adversarial search

CHAPTER 5 IN THE TEXTBOOK




Types of Games

- Many different kinds of games!
- Axes:
 - Deterministic or stochastic?
 - One, two, or more players?
 - Zero sum?
 - Perfect information (can you see the state)?
- Want algorithms for calculating a **strategy (policy)** which recommends a move from each state




Deterministic Games

- Many possible formalizations, one is:
 - States: S (start at s_0)
 - Players: $P = \{1 \dots N\}$ (usually take turns)
 - Actions: A (may depend on player / state)
 - Transition Function: $S \times A \rightarrow S$
 - Terminal Test: $S \rightarrow \{T, F\}$
 - Terminal Utilities: $S \times P \rightarrow \mathbb{R}$
- Solution for a player is a **policy**: $S \rightarrow A$



Zero-Sum Games



- Zero-Sum Games


- Agents have opposite utilities (values on outcomes)
- Lets us think of a single value that one maximizes and the other minimizes
- Adversarial, pure competition



- General Games


- Agents have independent utilities (values on outcomes)
- Cooperation, indifference, competition, and more are all possible
- More later on non-zero-sum games

Single-Agent Trees



Value of a State

Value of a state:
The best achievable
outcome (utility)
from that state




Non-Terminal States:

$$V(s) = \max_{s' \in \text{children}(s)} V(s')$$

Terminal States:

$$V(s) = \text{known}$$

Adversarial Game Trees




Minimax Values

States Under Agent's Control:

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$


States Under Opponent's Control:

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$



Terminal States:

$$V(s) = \text{known}$$



- In principle Chess is easy to solve:
 - A finite number of moves for both players to choose from
 - Build up a tree where the players take turns to choose piece movements on the board configuration that has evolved
 - No matter what search strategy we use, in practice all board configurations cannot be evaluated

Tic-Tac-Toe Game Tree



MAX (X)



MIN (O)





MAX (X)



MIN (O)


TERMINAL

Utility




Adversarial Search (Minimax)

- Deterministic, zero-sum games:
 - Tic-tac-toe, chess, checkers
 - One player maximizes result
 - The other minimizes result
- Minimax search:
 - A state-space search tree
 - Players alternate turns
 - Compute each node's **minimax value**: the best achievable utility against a rational (optimal) adversary



Optimal Decisions in Games

- The two players: min and max
- The initial board position is like the rules of the game dictate and max is the first to move
- Successor function $S(n)$ determines legal moves and resulting states
- A terminal test determines when the game is over
- max (min) aims at maximizing (minimizing) the value of the utility function
- The initial state and the successor function determine a *game tree*, where the players take turns to choose an edge to travel



- In our quest for the optimal game strategy, we will assume that also the adversary is infallible
- Player min chooses the moves that are best for it
- To determine the optimal strategy, we compute for each node n its *minimax* value $MM(n)$:

$$MM(n) = \begin{cases} \text{PAYOFF}(n), & \text{if } n \text{ is a terminal state} \\ \max_{s \in S(n)} MM(s), & \text{if } n \text{ is a max node} \\ \min_{s \in S(n)} MM(s), & \text{if } n \text{ is a min node} \end{cases}$$

Minimax Implementation


def max-value(state):

 initialize $v = -\infty$

 for each successor of state:

$v = \max(v, \text{min-value}(\text{successor}))$

 return v



def min-value(state):

 initialize $v = +\infty$

 for each successor of state:

$v = \min(v, \text{max-value}(\text{successor}))$

 return v


$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$


Minimax Implementation (Dispatch)

```
def value(state):  
    if the state is a terminal state: return the state's utility  
    if the next agent is MAX: return max-value(state)  
    if the next agent is MIN: return min-value(state)
```

```
def max-value(state):  
    initialize  $v = -\infty$   
    for each successor of state:  
         $v = \max(v, \text{value}(\text{successor}))$   
    return  $v$ 
```



```
def min-value(state):  
    initialize  $v = +\infty$   
    for each successor of state:  
         $v = \min(v, \text{value}(\text{successor}))$   
    return  $v$ 
```



- The play between two optimally playing players is completely determined by the minimax values
- For **max** the minimax values gives the worst-case outcome — the opponent **min** is optimal
- If the opponent does not choose the best moves, then **max** will do at least as well as against **min**
- Against suboptimal opponents there may be other strategies that do better than minimax
- The minimax algorithm performs a complete depth-first exploration of the game tree
- Therefore, the time complexity is $O(b^m)$, where b is the number of legal moves at each point and m is the maximum depth
- For real games, exponential time cost is totally impractical

Video of Demo Min vs. Rnd (Min)




Video of Demo Min vs. Rnd (Rnd)



Minimax Efficiency

- How efficient is minimax?
 - Just like (exhaustive) DFS
 - Time: $O(b^m)$
 - Space: $O(bm)$
- For chess, $b \approx 35$, $m \approx 100$
 - Exact solution is completely infeasible
 - But, do we need to explore the whole tree?
- Optimal against a perfect player. Otherwise?







- Examining the game tree we could find out which successive moves lead White to win and which of them lead to a win for Black
- The branching factor for chess is $b = 35$, and the length of longest play is infinite
- No matter what search strategy we use, in practice all board configurations cannot be evaluated
- Instead use a **payoff** (or utility) **function** to estimate how the board configuration evolves as moves are chosen
- The simplest payoff could be determined at the end of the game; did we win (1), draw ($\frac{1}{2}$), or lose (0)
- More applicable is to estimate any board position by, e.g., summing up the (difference of) material values of remaining pieces
 - pawn 1, knight 3, bishop 3, rook 5, queen 9

Resource Limits

- Problem: In realistic games, cannot search to leaves!
- Solution: Depth-limited search
 - Instead, search only to a limited depth in the tree
 - Replace terminal utilities with an evaluation function for non-terminal positions
- Example:
 - Suppose we have 100 seconds, can explore 10K nodes / sec
 - So can check 1M nodes per move
 - $\alpha - \beta$ reaches about depth 8 – decent chess program
- Guarantee of optimal play is gone
- More plies makes a BIG difference
- Use iterative deepening for an anytime algorithm



Video of Demo Limited Depth (2)





Video of Demo Limited Depth (10)



Alpha-beta pruning


- The exponential complexity of minimax search can be alleviated by pruning the nodes of the game tree that get evaluated
- It is possible to compute the correct minimax decision without looking at every node in the game tree
- Alpha-beta pruning gets its name from the parameters that describe bounds on the backed-up values that appear anywhere along the path
 - α = the value of the best (highest-value) choice we have found so far at any choice point along the path for **max**
 - β = the value of the best (lowest-value) choice we have found so far at any choice point along the path for **min**





Alpha-Beta Pruning

- General configuration (MIN version)
 - We're computing the MIN-VALUE at some node n
 - We're looping over n 's children
 - n 's estimate of the childrens' min is dropping
 - Who cares about n 's value? MAX
 - Let α be the best value that MAX can get at any choice point along the current path from the root
 - If n becomes worse than α , MAX will avoid it, so we can stop considering n 's other children (it's already bad enough that it won't be played)
- MAX version is symmetric




Alpha-Beta Implementation

α : MAX's best option on path to root
 β : MIN's best option on path to root

```
def max-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = -\infty$   
    for each successor of state:  
         $v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \geq \beta$  return  $v$   
         $\alpha = \max(\alpha, v)$   
    return  $v$ 
```

```
def min-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = +\infty$   
    for each successor of state:  
         $v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \leq \alpha$  return  $v$   
         $\beta = \min(\beta, v)$   
    return  $v$ 
```

Behavior from Computation



Video of Demo Mystery Pacman

